

---

## ▼ Capstone Project

### Image classifier for the SVHN dataset

#### Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course to build, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

#### How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

#### Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
import tensorflow as tf
from scipy.io import loadmat
```



SVHN overview image For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. "Reading Digits in Natural Images with Unsupervised Feature Learning". NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a world image into one of ten classes.

```
# Run this cell to connect to your Drive folder
```

```
from google.colab import drive
drive.mount('/content/gdrive')
```

➞ Go to this URL in a browser: [https://accounts.google.com/o/oauth2/auth?client\\_id=947318989803-6bn6qk8qdgf4n4g](https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g)

```
Enter your authorization code:
.....
Mounted at /content/gdrive
```

```
# Run this cell to load the dataset
```

```
# train = loadmat('data/train_32x32.mat')
# test = loadmat('data/test_32x32.mat')
```

```
# Load the dataset from your Drive folder
```

```
train = loadmat('/content/gdrive/My Drive/TensorFlow/train_32x32.mat')
test = loadmat('/content/gdrive/My Drive/TensorFlow/test_32x32.mat')
```

Both `train` and `test` are dictionaries with keys `x` and `y` for the input images and labels respectively.

## ▼ 1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.

- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

```
X_train , y_train = train['X'] , train['y']
X_test , y_test = test['X'] , test['y']
print(X_train.shape)
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

```
↳ (32, 32, 3, 73257)
   (73257, 1)
   (32, 32, 3, 26032)
   (26032, 1)
```

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.io import loadmat
from skimage import color
from skimage import io
from sklearn.model_selection import train_test_split
```

```
%matplotlib inline
plt.rcParams['figure.figsize'] = (16.0, 4.0)
```

```
↳ /usr/local/lib/python3.6/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning: pandas.util.testing is deprecated
   import pandas.util.testing as tm
```

```
X_train, y_train = X_train.transpose((3,0,1,2)), y_train[:,0]
X_test, y_test = X_test.transpose((3,0,1,2)), y_test[:,0]
```

```
print("Training Set", X_train.shape)
print("Test Set", X_test.shape)
print('')
```

```
# Calculate the total number of images
num_images = X_train.shape[0] + X_test.shape[0]

print("Total Number of Images", num_images)
```

```
↳ Training Set (73257, 32, 32, 3)
   Test Set (26032, 32, 32, 3)
```

```
Total Number of Images 99289
```

```
def plot_images(img, labels, nrows, ncols):
    """ Plot nrows x ncols images
    """
    fig, axes = plt.subplots(nrows, ncols)
    for i, ax in enumerate(axes.flat):
        if img[i].shape == (32, 32, 3):
            ax.imshow(img[i])
        else:
            ax.imshow(img[i,:,:,:0])
        ax.set_xticks([]); ax.set_yticks([])
        ax.set_title(labels[i])
```

```
# Plot some training set images
plot_images(X_train, y_train, 2, 8)
```

```
↳
```



```
# Plot some test set images
plot_images(X_test, y_test, 2, 8)
```



```
fig, (ax1, ax2) = plt.subplots(1, 2, sharex=True)

fig.suptitle('Class Distribution', fontsize=14, fontweight='bold', y=1.05)


ax1.hist(y_train, bins=10)
ax1.set_title("Training set")
ax1.set_xlim(1, 10)

ax2.hist(y_test, color='g', bins=10)
ax2.set_title("Test set")

fig.tight_layout()
```

↪

## Class Distribution



```
# Converting Label 10 -> 0
y_train[y_train == 10] = 0
y_test[y_test == 10] = 0
```

Splitting the Training to Train+Validation Splitting to 13% in Val Set as it gives around 9500 data having min. of 800 instances of each class

Using random state to regenerate the whole Dataset in re-run

```
X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.13, random_state=7)
```

Visualize New Distribution

```
fig, (ax1, ax2) = plt.subplots(1, 2, sharex=True)

fig.suptitle('Class Distribution', fontsize=14, fontweight='bold', y=1.05)


ax1.hist(y_train, bins=10)
ax1.set_title("Training set")
ax1.set_xlim(1, 10)
```

```
ax2.hist(y_val, color='g', bins=10)
ax2.set_title("Validation set")

fig.tight_layout()
```



**Class Distribution**



```
y_train.shape, y_val.shape, y_test.shape
```

```
((63733,), (9524,), (26032,))
```

## Grayscale Conversion

To speed up our experiments we will convert our images from RGB to Grayscale, which greatly reduces the amount of data we will have to process.

$$Y = 0.2990R + 0.5870G + 0.1140B$$

Here is a simple function that helps us print the size of a numpy array in a human readable format.

```
def rgb2gray(images):
    return np.expand_dims(np.dot(images, [0.2990, 0.5870, 0.1140]), axis=3)
```

```
return np.expand_dims(np.dot(images, [0.299, 0.587, 0.114]), axis=3)
```

## Converting to float for numpy computation


```
train_greyscale = rgb2gray(X_train).astype(np.float32)
test_greyscale = rgb2gray(X_test).astype(np.float32)
val_greyscale = rgb2gray(X_val).astype(np.float32)
```

```
print("Training Set", train_greyscale.shape)
print("Validation Set", val_greyscale.shape)
print("Test Set", test_greyscale.shape)
print('')
```

```
↳ Training Set (63733, 32, 32, 1)
   Validation Set (9524, 32, 32, 1)
   Test Set (26032, 32, 32, 1)
```

```
del X_train, X_test, X_val
```

```
plot_images(train_greyscale, y_train, 1, 10)
```



## Doing Normalization



```

# Calculate the mean on the training data
train_mean = np.mean(train_greyscale, axis=0)

# Calculate the std on the training data
train_std = np.std(train_greyscale, axis=0)

# Subtract it equally from all splits
train_greyscale_norm = (train_greyscale - train_mean) / train_std
test_greyscale_norm = (test_greyscale - train_mean) / train_std
val_greyscale_norm = (val_greyscale - train_mean) / train_std

from sklearn.preprocessing import OneHotEncoder

# Fit the OneHotEncoder
enc = OneHotEncoder().fit(y_train.reshape(-1, 1))

# Transform the label values to a one-hot-encoding scheme
y_train = enc.transform(y_train.reshape(-1, 1)).toarray()
y_test = enc.transform(y_test.reshape(-1, 1)).toarray()
y_val = enc.transform(y_val.reshape(-1, 1)).toarray()

print("Training set", y_train.shape)
print("Validation set", y_val.shape)
print("Test set", y_test.shape)

```

```

☞ Training set (63733, 10)
   Validation set (9524, 10)
   Test set (26032, 10)

```

## Storing Data to Disk

Stored only the Grayscale Data not the RGB

```

import h5py

# Create file

```

```

# Create h5f
h5f = h5py.File('SVHN_grey.h5', 'w')

# Store the datasets
h5f.create_dataset('X_train', data=train_greyscale_norm)
h5f.create_dataset('y_train', data=y_train)
h5f.create_dataset('X_test', data=test_greyscale_norm)
h5f.create_dataset('y_test', data=y_test)
h5f.create_dataset('X_val', data=val_greyscale_norm)
h5f.create_dataset('y_val', data=y_val)

# Close the file
h5f.close()

```

## ▼ 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer h 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reason accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the summary() method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during th training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be h
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```

import os
import time
# from __future__ import absolute_import
# from __future__ import print_function

```

```
from datetime import timedelta
import h5py
import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
plt.rcParams['figure.figsize'] = (16.0, 4.0) # Set default figure size
```

```
h5f = h5py.File('SVHN_grey.h5', 'r')
```

```
# Load the training, test and validation set
```

```
X_train = h5f['X_train'][:]
```

```
y_train = h5f['y_train'][:]
```

```
X_test = h5f['X_test'][:]
```

```
y_test = h5f['y_test'][:]
```

```
X_val = h5f['X_val'][:]
```

```
y_val = h5f['y_val'][:]
```

```
# Close this file
```

```
h5f.close()
```

```
print('Training set', X_train.shape, y_train.shape)
```

```
print('Validation set', X_val.shape, y_val.shape)
```

```
print('Test set', X_test.shape, y_test.shape)
```

```
↳ Training set (63733, 32, 32, 1) (63733, 10)
   Validation set (9524, 32, 32, 1) (9524, 10)
   Test set (26032, 32, 32, 1) (26032, 10)
```

```
# Display one of the images
```


```
i = 0
```

```
labels = np.argmax(y_train[i])
```

```
img = X_train[i, :, :, 0]
```


```
plt.imshow(img)
```

```
plt.show()
print(f"label: {labels}")
```



```
X_train[0].shape
import matplotlib.pyplot as plt
```

```
fig, ax = plt.subplots(1, 10, figsize=(10, 1))
for i in range(10):
    ax[i].set_axis_off()
    ax[i].imshow(X_train[i, :, :, 0])
```



```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
```

```
def get_model(input_shape):

    model = Sequential([
```

```

    Flatten(input_shape = input_shape),
    Dense(128, activation = 'relu'),
    Dense(128, activation = 'relu'),
    Dense(128, activation = 'relu'),
    Dense(128, activation = 'relu'),
    Dense(128, activation = 'relu'),
    Dense(10, activation = 'softmax')
])

return model

```

```

model = get_model(X_train[0].shape)
model.summary()

```

☞ Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 128)	131200
dense_1 (Dense)	(None, 128)	16512
dense_2 (Dense)	(None, 128)	16512
dense_3 (Dense)	(None, 128)	16512
dense_4 (Dense)	(None, 128)	16512
dense_5 (Dense)	(None, 10)	1290
=====		
Total params: 198,538		
Trainable params: 198,538		
Non-trainable params: 0		

```

def compile_model(model):

```

```

model.compile(
    optimizer = tf.keras.optimizers.Adam(learning_rate=0.0001),
    loss = 'categorical_crossentropy',
    metrics = ['accuracy']
)

```

```

compile_model(model)
print(model.optimizer)
print(model.loss)
print(model.metrics)
print(model.optimizer.lr)

```

```

❏ <tensorflow.python.keras.optimizer_v2.adam.Adam object at 0x7f4e80342be0>
categorical_crossentropy
[]
<tf.Variable 'learning_rate:0' shape=() dtype=float32, numpy=1e-04>

```

```

from tensorflow.keras.callbacks import Callback, ModelCheckpoint

```

```

class TrainingCallback(Callback):

```

```

    def on_train_begin(self, logs=None):
        print("Starting training....")

```

```

    def on_epoch_begin(self, epoch, logs=None):
        print(f"Starting epoch {epoch}")

```

```

    def on_epoch_end(self, epoch, logs=None):
        print(f"Finishing epoch {epoch}")

```

```

    def on_train_end(self, logs=None):
        print("Finished training:")

```

```

def get_checkpoint_best_only():

```

```

def get_checkpoint_best_only():

    checkpoint_best_path = 'checkpoints_best_only/checkpoint'
    checkpoint_best_only = ModelCheckpoint(filepath=checkpoint_best_path, save_freq='epoch',
                                           save_weights_only=True, monitor = 'val_accuracy',
                                           save_best_only=True, verbose = 1)

    return checkpoint_best_only


TrainingCallback = TrainingCallback()
checkpoint_best_only = get_checkpoint_best_only()


def train_model(model, train_data, train_targets, epochs):

    history = model.fit(train_data, train_targets, epochs=epochs,
                       batch_size=64, validation_data=(X_val,y_val),verbose=False)

    return history


callbacks = [TrainingCallback, checkpoint_best_only]
history = model.fit(X_train, y_train, epochs=30, batch_size=64, validation_data=(X_val,y_val),callbacks=callbacks)

```



Starting training....

Starting epoch 0

Epoch 1/30

985/996 [=====>.] - ETA: 0s - loss: 1.5683 - accuracy: 0.4728Finishing epoch 0

Epoch 00001: val\_accuracy improved from -inf to 0.65981, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 1.5620 - accuracy: 0.4753 - val\_loss: 1.0931 -

Starting epoch 1

Epoch 2/30

987/996 [=====>.] - ETA: 0s - loss: 0.9865 - accuracy: 0.6966Finishing epoch 1

Epoch 00002: val\_accuracy improved from 0.65981 to 0.72764, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.9857 - accuracy: 0.6968 - val\_loss: 0.9042 -

Starting epoch 2

Epoch 3/30

991/996 [=====>.] - ETA: 0s - loss: 0.8452 - accuracy: 0.7430Finishing epoch 2

Epoch 00003: val\_accuracy improved from 0.72764 to 0.75430, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.8449 - accuracy: 0.7431 - val\_loss: 0.8217 -

Starting epoch 3

Epoch 4/30

986/996 [=====>.] - ETA: 0s - loss: 0.7582 - accuracy: 0.7699Finishing epoch 3

Epoch 00004: val\_accuracy improved from 0.75430 to 0.77572, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.7575 - accuracy: 0.7702 - val\_loss: 0.7477 -

Starting epoch 4

Epoch 5/30

990/996 [=====>.] - ETA: 0s - loss: 0.6957 - accuracy: 0.7895Finishing epoch 4

Epoch 00005: val\_accuracy improved from 0.77572 to 0.79074, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.6954 - accuracy: 0.7894 - val\_loss: 0.7049 -

Starting epoch 5

Epoch 6/30

985/996 [=====>.] - ETA: 0s - loss: 0.6499 - accuracy: 0.8030Finishing epoch 5

Epoch 00006: val\_accuracy improved from 0.79074 to 0.79473, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.6500 - accuracy: 0.8029 - val\_loss: 0.6912 -

Starting epoch 6

Epoch 7/30

989/996 [=====>.] - ETA: 0s - loss: 0.6133 - accuracy: 0.8163Finishing epoch 6

Epoch 00007: val\_accuracy improved from 0.79473 to 0.80292, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.6140 - accuracy: 0.8161 - val\_loss: 0.6678 -



Starting epoch 7

Epoch 8/30

994/996 [=====>.] - ETA: 0s - loss: 0.5829 - accuracy: 0.8248Finishing epoch 7

Epoch 00008: val\_accuracy improved from 0.80292 to 0.80617, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.5832 - accuracy: 0.8247 - val\_loss: 0.6542 -

Starting epoch 8

Epoch 9/30

989/996 [=====>.] - ETA: 0s - loss: 0.5560 - accuracy: 0.8332Finishing epoch 8

Epoch 00009: val\_accuracy improved from 0.80617 to 0.81142, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.5556 - accuracy: 0.8335 - val\_loss: 0.6367 -

Starting epoch 9

Epoch 10/30

992/996 [=====>.] - ETA: 0s - loss: 0.5326 - accuracy: 0.8380Finishing epoch 9

Epoch 00010: val\_accuracy improved from 0.81142 to 0.81804, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.5328 - accuracy: 0.8380 - val\_loss: 0.6250 -

Starting epoch 10

Epoch 11/30

989/996 [=====>.] - ETA: 0s - loss: 0.5144 - accuracy: 0.8444Finishing epoch 10

Epoch 00011: val\_accuracy did not improve from 0.81804

996/996 [=====] - 5s 5ms/step - loss: 0.5141 - accuracy: 0.8445 - val\_loss: 0.6171 -

Starting epoch 11

Epoch 12/30

992/996 [=====>.] - ETA: 0s - loss: 0.4970 - accuracy: 0.8506Finishing epoch 11

Epoch 00012: val\_accuracy improved from 0.81804 to 0.82392, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.4969 - accuracy: 0.8506 - val\_loss: 0.6131 -

Starting epoch 12

Epoch 13/30

986/996 [=====>.] - ETA: 0s - loss: 0.4788 - accuracy: 0.8552Finishing epoch 12

Epoch 00013: val\_accuracy improved from 0.82392 to 0.82812, saving model to checkpoints\_best\_only/checkpoint

996/996 [=====] - 5s 5ms/step - loss: 0.4789 - accuracy: 0.8552 - val\_loss: 0.5965 -

Starting epoch 13

Epoch 14/30

991/996 [=====>.] - ETA: 0s - loss: 0.4663 - accuracy: 0.8584Finishing epoch 13

Epoch 00014: val\_accuracy did not improve from 0.82812

996/996 [=====] - 5s 5ms/step - loss: 0.4666 - accuracy: 0.8583 - val\_loss: 0.5997 -

Starting epoch 14

Epoch 15/30

991/996 [=====>.] - ETA: 0s - loss: 0.4492 - accuracy: 0.8634Finishing epoch 14

Epoch 00015: val\_accuracy improved from 0.82812 to 0.82980, saving model to checkpoints\_best\_only/checkpoint  
996/996 [=====] - 5s 5ms/step - loss: 0.4494 - accuracy: 0.8633 - val\_loss: 0.5943 -  
Starting epoch 15  
Epoch 16/30  
993/996 [=====>.] - ETA: 0s - loss: 0.4393 - accuracy: 0.8656Finishing epoch 15

Epoch 00016: val\_accuracy did not improve from 0.82980  
996/996 [=====] - 5s 5ms/step - loss: 0.4395 - accuracy: 0.8656 - val\_loss: 0.5958 -  
Starting epoch 16  
Epoch 17/30  
991/996 [=====>.] - ETA: 0s - loss: 0.4241 - accuracy: 0.8706Finishing epoch 16

Epoch 00017: val\_accuracy improved from 0.82980 to 0.83358, saving model to checkpoints\_best\_only/checkpoint  
996/996 [=====] - 5s 5ms/step - loss: 0.4239 - accuracy: 0.8706 - val\_loss: 0.5829 -  
Starting epoch 17  
Epoch 18/30  
988/996 [=====>.] - ETA: 0s - loss: 0.4148 - accuracy: 0.8748Finishing epoch 17

Epoch 00018: val\_accuracy did not improve from 0.83358  
996/996 [=====] - 5s 5ms/step - loss: 0.4150 - accuracy: 0.8747 - val\_loss: 0.5880 -  
Starting epoch 18  
Epoch 19/30  
986/996 [=====>.] - ETA: 0s - loss: 0.4028 - accuracy: 0.8775Finishing epoch 18

Epoch 00019: val\_accuracy improved from 0.83358 to 0.83830, saving model to checkpoints\_best\_only/checkpoint  
996/996 [=====] - 5s 5ms/step - loss: 0.4027 - accuracy: 0.8775 - val\_loss: 0.5709 -  
Starting epoch 19  
Epoch 20/30  
992/996 [=====>.] - ETA: 0s - loss: 0.3928 - accuracy: 0.8805Finishing epoch 19

Epoch 00020: val\_accuracy did not improve from 0.83830  
996/996 [=====] - 5s 5ms/step - loss: 0.3931 - accuracy: 0.8804 - val\_loss: 0.5877 -  
Starting epoch 20  
Epoch 21/30  
994/996 [=====>.] - ETA: 0s - loss: 0.3849 - accuracy: 0.8816Finishing epoch 20

Epoch 00021: val\_accuracy did not improve from 0.83830  
996/996 [=====] - 5s 5ms/step - loss: 0.3848 - accuracy: 0.8816 - val\_loss: 0.5958 -  
Starting epoch 21  
Epoch 22/30  
992/996 [=====>.] - ETA: 0s - loss: 0.3758 - accuracy: 0.8846Finishing epoch 21

Epoch 00022: val\_accuracy did not improve from 0.83830  
 996/996 [=====] - 5s 5ms/step - loss: 0.3756 - accuracy: 0.8847 - val\_loss: 0.5969 -  
 Starting epoch 22  
 Epoch 23/30  
 993/996 [=====>.] - ETA: 0s - loss: 0.3656 - accuracy: 0.8878Finishing epoch 22

Epoch 00023: val\_accuracy improved from 0.83830 to 0.83956, saving model to checkpoints\_best\_only/checkpoint  
 996/996 [=====] - 5s 5ms/step - loss: 0.3654 - accuracy: 0.8878 - val\_loss: 0.5798 -  
 Starting epoch 23  
 Epoch 24/30  
 996/996 [=====] - ETA: 0s - loss: 0.3579 - accuracy: 0.8900Finishing epoch 23

Epoch 00024: val\_accuracy did not improve from 0.83956  
 996/996 [=====] - 5s 5ms/step - loss: 0.3579 - accuracy: 0.8900 - val\_loss: 0.6121 -  
 Starting epoch 24  
 Epoch 25/30  
 995/996 [=====>.] - ETA: 0s - loss: 0.3488 - accuracy: 0.8930Finishing epoch 24

Epoch 00025: val\_accuracy improved from 0.83956 to 0.84313, saving model to checkpoints\_best\_only/checkpoint  
 996/996 [=====] - 5s 5ms/step - loss: 0.3486 - accuracy: 0.8930 - val\_loss: 0.5813 -  
 Starting epoch 25  
 Epoch 26/30  
 988/996 [=====>.] - ETA: 0s - loss: 0.3402 - accuracy: 0.8962Finishing epoch 25

Epoch 00026: val\_accuracy did not improve from 0.84313  
 996/996 [=====] - 5s 5ms/step - loss: 0.3407 - accuracy: 0.8960 - val\_loss: 0.5783 -  
 Starting epoch 26  
 Epoch 27/30  
 991/996 [=====>.] - ETA: 0s - loss: 0.3314 - accuracy: 0.8987Finishing epoch 26

Epoch 00027: val\_accuracy did not improve from 0.84313  
 996/996 [=====] - 5s 5ms/step - loss: 0.3320 - accuracy: 0.8985 - val\_loss: 0.5780 -  
 Starting epoch 27  
 Epoch 28/30  
 996/996 [=====] - ETA: 0s - loss: 0.3241 - accuracy: 0.9009Finishing epoch 27

Epoch 00028: val\_accuracy did not improve from 0.84313  
 996/996 [=====] - 5s 5ms/step - loss: 0.3241 - accuracy: 0.9009 - val\_loss: 0.5866 -  
 Starting epoch 28


try:

```
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
```

```

except KeyError:
    plt.plot(history.history['acc'])
    plt.plot(history.history['val_acc'])
plt.title('Accuracy vs. epochs')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='lower right')
plt.show()

```




```

#Run this cell to plot the epoch vs loss graph
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()

```





```
import pandas as pd
df = pd.DataFrame(history.history)
df.head(30)
```



- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
h5f = h5py.File('SVHN_grey.h5', 'r')

# Load the training, test and validation set
X_train = h5f['X_train'][:]
y_train = h5f['y_train'][:]
X_test = h5f['X_test'][:]
y_test = h5f['y_test'][:]
X_val = h5f['X_val'][:]
y_val = h5f['y_val'][:]

# Close this file
h5f.close()

print('Training set', X_train.shape, y_train.shape)
print('Validation set', X_val.shape, y_val.shape)
print('Test set', X_test.shape, y_test.shape)
```



	<b>loss</b>	<b>accuracy</b>	<b>val_loss</b>	<b>val_accuracy</b>
<b>0</b>	1.561980	0.475327	1.093143	0.659807
<b>1</b>	0.985651	0.696813	0.904183	0.727635
<b>2</b>	0.844937	0.743053	0.821673	0.754305
<b>3</b>	0.757547	0.770166	0.747658	0.775724
<b>4</b>	0.695384	0.789434	0.704870	0.790739
<b>5</b>	0.649967	0.802881	0.691177	0.794729
<b>6</b>	0.613982	0.816139	0.667764	0.802919
<b>7</b>	0.583163	0.824738	0.654172	0.806174
<b>8</b>	0.555633	0.833477	0.636668	0.811424
<b>9</b>	0.532793	0.837965	0.625032	0.818039
<b>10</b>	0.514102	0.844460	0.617072	0.817829
<b>11</b>	0.496858	0.850595	0.613067	0.823919
<b>12</b>	0.478938	0.855161	0.596519	0.828118
<b>13</b>	0.466632	0.858315	0.599689	0.827383
<b>14</b>	0.449364	0.863336	0.594293	0.829798
<b>15</b>	0.439456	0.865627	0.595806	0.829588
<b>16</b>	0.423869	0.870601	0.582872	0.833578

```
test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=2)
```

```
814/814 - 2s - loss: 0.7499 - accuracy: 0.8161
```

```
0.423869 0.870601 0.582872 0.833578
```


### ▼ 3. CNN neural network classifier

```

Training set (63733, 32, 32, 1) (63733, 10)

# Display one of the images
i = 30
labels = np.argmax(y_train[i])
img = X_train[i,:,:,0]
plt.imshow(img)
plt.show()
print(f"label: {labels}")

```




```

import matplotlib.pyplot as plt

fig, ax = plt.subplots(1, 10, figsize=(10, 1))
for i in range(10):
    ax[i].set_axis_off()
    ax[i].imshow(X_train[i,:,:,0])

```





```

import tensorflow as tf
from tensorflow.keras.preprocessing.image import load_img, img_to_array
from tensorflow.keras.models import Sequential, load_model
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, BatchNormalization, Dropout
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
import os
import numpy as np
import pandas as pd

```

```

def get_new_model(input_shape):

```

```

    model = Sequential([
        Conv2D(16,(3,3),padding="SAME",activation='relu',name='conv_1', input_shape=(input_shape)),
        Dropout(0.5),
        Conv2D(8,(3,3),padding="SAME",activation='relu',name='conv_2'),
        BatchNormalization(),
        Dropout(0.5),
        MaxPooling2D((4,4),name='pool_1'),
        Dense(128,activation='relu',name='dense_1'),
        Flatten(name='flatten2'),
        Dense(10,activation='softmax',name='dense_2')
    ])

```

```

    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

```

```

    return model

```

```

model = get_new_model(X_train[0].shape)
model.summary()

```



Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv_1 (Conv2D)	(None, 32, 32, 16)	160
dropout (Dropout)	(None, 32, 32, 16)	0
conv_2 (Conv2D)	(None, 32, 32, 8)	1160
batch_normalization (Batch Normalization)	(None, 32, 32, 8)	32
dropout_1 (Dropout)	(None, 32, 32, 8)	0
pool_1 (MaxPooling2D)	(None, 8, 8, 8)	0
dense_1 (Dense)	(None, 8, 8, 128)	1152
flatten2 (Flatten)	(None, 8192)	0
dense_2 (Dense)	(None, 10)	81930
Total params: 84,434		
Trainable params: 84,418		

```
compile_model(model)
print(model.optimizer)
print(model.loss)
print(model.metrics)
print(model.optimizer.lr)
```

```
[<tensorflow.python.keras.optimizer_v2.adam.Adam object at 0x7f4e8008f7f0>
categorical_crossentropy
[]
<tf.Variable 'learning_rate:0' shape=() dtype=float32, numpy=1e-04>
```

```
def get_early_stopping():

    early_stopping = EarlyStopping(monitor='val_accuracy', patience = 7)

    return early_stopping
```

```
return early_stopping
```

```
def get_checkpoint_best_CNN():
```

```
    checkpoint_best_path = 'checkpoints_best_only_CNN/checkpoint'
```

```
    checkpoint_best_only_CNN = ModelCheckpoint(filepath=checkpoint_best_path, save_freq='epoch',  
                                                save_weights_only=True, monitor = 'val_accuracy',  
                                                save_best_only=True, verbose = 1)
```

```
    return checkpoint_best_only_CNN
```

```
early_stopping = get_early_stopping()
```

```
checkpoint_best_only_CNN = get_checkpoint_best_CNN()
```

```
callbacks = [checkpoint_best_only_CNN, early_stopping]
```

```
history = model.fit(X_train, y_train, epochs=30, batch_size=64, validation_data=(X_val,y_val),callbacks=callbacks)
```



Epoch 1/30  
996/996 [=====] - ETA: 0s - loss: 1.9911 - accuracy: 0.3186  
Epoch 00001: val\_accuracy improved from -inf to 0.47848, saving model to checkpoints\_best\_only\_CNN/checkpoint  
996/996 [=====] - 10s 10ms/step - loss: 1.9911 - accuracy: 0.3186 - val\_loss: 1.5333  
Epoch 2/30  
991/996 [=====>.] - ETA: 0s - loss: 1.0562 - accuracy: 0.6781  
Epoch 00002: val\_accuracy improved from 0.47848 to 0.68490, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 1.0554 - accuracy: 0.6784 - val\_loss: 1.0023 -  
Epoch 3/30  
990/996 [=====>.] - ETA: 0s - loss: 0.8436 - accuracy: 0.7510  
Epoch 00003: val\_accuracy improved from 0.68490 to 0.72071, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 0.8428 - accuracy: 0.7511 - val\_loss: 0.8803 -  
Epoch 4/30  
995/996 [=====>.] - ETA: 0s - loss: 0.7528 - accuracy: 0.7814  
Epoch 00004: val\_accuracy improved from 0.72071 to 0.75231, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 0.7527 - accuracy: 0.7814 - val\_loss: 0.8121 -  
Epoch 5/30  
996/996 [=====] - ETA: 0s - loss: 0.7008 - accuracy: 0.7978  
Epoch 00005: val\_accuracy improved from 0.75231 to 0.75493, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 0.7008 - accuracy: 0.7978 - val\_loss: 0.7873 -  
Epoch 6/30  
991/996 [=====>.] - ETA: 0s - loss: 0.6637 - accuracy: 0.8082  
Epoch 00006: val\_accuracy improved from 0.75493 to 0.77110, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 0.6631 - accuracy: 0.8082 - val\_loss: 0.7438 -  
Epoch 7/30  
993/996 [=====>.] - ETA: 0s - loss: 0.6322 - accuracy: 0.8170  
Epoch 00007: val\_accuracy improved from 0.77110 to 0.78402, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 0.6322 - accuracy: 0.8170 - val\_loss: 0.7134 -  
Epoch 8/30  
993/996 [=====>.] - ETA: 0s - loss: 0.6080 - accuracy: 0.8231  
Epoch 00008: val\_accuracy improved from 0.78402 to 0.79242, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 0.6080 - accuracy: 0.8232 - val\_loss: 0.6862 -  
Epoch 9/30  
991/996 [=====>.] - ETA: 0s - loss: 0.5915 - accuracy: 0.8286  
Epoch 00009: val\_accuracy did not improve from 0.79242  
996/996 [=====] - 9s 9ms/step - loss: 0.5910 - accuracy: 0.8288 - val\_loss: 0.6878 -  
Epoch 10/30  
995/996 [=====>.] - ETA: 0s - loss: 0.5739 - accuracy: 0.8332  
Epoch 00010: val\_accuracy improved from 0.79242 to 0.79966, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 0.5739 - accuracy: 0.8332 - val\_loss: 0.6625 -  
Epoch 11/30  
995/996 [=====>.] - ETA: 0s - loss: 0.5585 - accuracy: 0.8368  
Epoch 00011: val\_accuracy did not improve from 0.79966

996/996 [=====] - 9s 9ms/step - loss: 0.5585 - accuracy: 0.8368 - val\_loss: 0.6626 -  
Epoch 12/30  
992/996 [=====>.] - ETA: 0s - loss: 0.5474 - accuracy: 0.8405  
Epoch 00012: val\_accuracy did not improve from 0.79966  
996/996 [=====] - 9s 9ms/step - loss: 0.5473 - accuracy: 0.8405 - val\_loss: 0.6556 -  
Epoch 13/30  
992/996 [=====>.] - ETA: 0s - loss: 0.5348 - accuracy: 0.8422  
Epoch 00013: val\_accuracy did not improve from 0.79966  
996/996 [=====] - 9s 9ms/step - loss: 0.5348 - accuracy: 0.8422 - val\_loss: 0.6534 -  
Epoch 14/30  
992/996 [=====>.] - ETA: 0s - loss: 0.5263 - accuracy: 0.8462  
Epoch 00014: val\_accuracy improved from 0.79966 to 0.80250, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 0.5265 - accuracy: 0.8461 - val\_loss: 0.6410 -  
Epoch 15/30  
992/996 [=====>.] - ETA: 0s - loss: 0.5202 - accuracy: 0.8457  
Epoch 00015: val\_accuracy improved from 0.80250 to 0.81562, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 0.5201 - accuracy: 0.8457 - val\_loss: 0.6129 -  
Epoch 16/30  
994/996 [=====>.] - ETA: 0s - loss: 0.5110 - accuracy: 0.8497  
Epoch 00016: val\_accuracy improved from 0.81562 to 0.82182, saving model to checkpoints\_best\_only\_CNN/checkpo:  
996/996 [=====] - 9s 9ms/step - loss: 0.5109 - accuracy: 0.8497 - val\_loss: 0.5984 -  
Epoch 17/30  
993/996 [=====>.] - ETA: 0s - loss: 0.5052 - accuracy: 0.8504  
Epoch 00017: val\_accuracy did not improve from 0.82182  
996/996 [=====] - 9s 9ms/step - loss: 0.5051 - accuracy: 0.8504 - val\_loss: 0.5956 -  
Epoch 18/30  
992/996 [=====>.] - ETA: 0s - loss: 0.4953 - accuracy: 0.8533  
Epoch 00018: val\_accuracy did not improve from 0.82182  
996/996 [=====] - 9s 9ms/step - loss: 0.4954 - accuracy: 0.8533 - val\_loss: 0.5999 -  
Epoch 19/30  
994/996 [=====>.] - ETA: 0s - loss: 0.4914 - accuracy: 0.8548  
Epoch 00019: val\_accuracy did not improve from 0.82182  
996/996 [=====] - 9s 9ms/step - loss: 0.4915 - accuracy: 0.8548 - val\_loss: 0.6015 -  
Epoch 20/30  
991/996 [=====>.] - ETA: 0s - loss: 0.4837 - accuracy: 0.8567  
Epoch 00020: val\_accuracy did not improve from 0.82182  
996/996 [=====] - 9s 9ms/step - loss: 0.4837 - accuracy: 0.8566 - val\_loss: 0.5930 -  
Epoch 21/30  
991/996 [=====>.] - ETA: 0s - loss: 0.4794 - accuracy: 0.8577  
Epoch 00021: val\_accuracy did not improve from 0.82182  
996/996 [=====] - 9s 9ms/step - loss: 0.4794 - accuracy: 0.8577 - val\_loss: 0.5948 -  
Epoch 22/30  
994/996 [=====>.] - ETA: 0s - loss: 0.4760 - accuracy: 0.8592  
Epoch 00022: val accuracy improved from 0.82182 to 0.82455, saving model to checkpoints best only CNN/checkpo:

```

996/996 [=====] - 9s 9ms/step - loss: 0.4759 - accuracy: 0.8593 - val_loss: 0.5777 -
Epoch 23/30
993/996 [=====>.] - ETA: 0s - loss: 0.4690 - accuracy: 0.8598
Epoch 00023: val_accuracy did not improve from 0.82455
996/996 [=====] - 9s 9ms/step - loss: 0.4690 - accuracy: 0.8598 - val_loss: 0.5930 -
Epoch 24/30
993/996 [=====>.] - ETA: 0s - loss: 0.4662 - accuracy: 0.8618
Epoch 00024: val_accuracy improved from 0.82455 to 0.82969, saving model to checkpoints_best_only_CNN/checkpo:
996/996 [=====] - 9s 9ms/step - loss: 0.4662 - accuracy: 0.8618 - val_loss: 0.5667 -
Epoch 25/30
996/996 [=====] - ETA: 0s - loss: 0.4638 - accuracy: 0.8622
Epoch 00025: val_accuracy did not improve from 0.82969
996/996 [=====] - 10s 10ms/step - loss: 0.4638 - accuracy: 0.8622 - val_loss: 0.5959
Epoch 26/30
996/996 [=====] - ETA: 0s - loss: 0.4587 - accuracy: 0.8628
Epoch 00026: val_accuracy improved from 0.82969 to 0.83379, saving model to checkpoints_best_only_CNN/checkpo:
996/996 [=====] - 9s 9ms/step - loss: 0.4587 - accuracy: 0.8628 - val_loss: 0.5549 -

```

try:

```

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])

```

except KeyError:

```

plt.plot(history.history['acc'])
plt.plot(history.history['val_acc'])

```


```

plt.title('Accuracy vs. epochs')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='lower right')
plt.show()

```




Accuracy vs. epochs



```
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')
plt.show()
```



Loss vs. epochs



```
import pandas as pd
df = pd.DataFrame(history.history)
df.head(30)
```



	<b>loss</b>	<b>accuracy</b>	<b>val_loss</b>	<b>val_accuracy</b>
<b>0</b>	1.991114	0.318563	1.533335	0.478475
<b>1</b>	1.055429	0.678361	1.002318	0.684901
<b>2</b>	0.842808	0.751134	0.880338	0.720706
<b>3</b>	0.752694	0.781432	0.812097	0.752310
<b>4</b>	0.700818	0.797781	0.787272	0.754935
<b>5</b>	0.663128	0.808247	0.743797	0.771105
<b>6</b>	0.632209	0.817002	0.713448	0.784019
<b>7</b>	0.607964	0.823153	0.686152	0.792419
<b>8</b>	0.591037	0.828754	0.687796	0.790109
<b>9</b>	0.573904	0.833179	0.662540	0.799664
<b>10</b>	0.558532	0.836788	0.662647	0.796304
<b>11</b>	0.547331	0.840522	0.655577	0.797459
<b>12</b>	0.534765	0.842170	0.653356	0.796724
<b>13</b>	0.526503	0.846124	0.641011	0.802499
<b>14</b>	0.520106	0.845731	0.612947	0.815624
<b>15</b>	0.510939	0.849717	0.598359	0.821819
<b>16</b>	0.505103	0.850423	0.595601	0.819929
<b>17</b>	0.495357	0.853294	0.599889	0.818564
<b>18</b>	0.491547	0.854753	0.601457	0.818354
<b>19</b>	0.483743	0.856589	0.593048	0.819089
<b>20</b>	0.479432	0.857672	0.594792	0.817829
<b>21</b>	0.475936	0.859272	0.577689	0.824548
<b>22</b>	0.468976	0.859774	0.592951	0.816884



<b>23</b>	0.466246	0.861845	0.566677	0.829693
<b>24</b>	0.463752	0.862191	0.595900	0.814574
<b>25</b>	0.458705	0.862771	0.554933	0.833788
<b>26</b>	0.458083	0.864074	0.561464	0.830428
<b>27</b>	0.455302	0.864623	0.561529	0.829798
<b>28</b>	0.452050	0.865203	0.560792	0.828748

```
test_loss, test_accuracy = model.evaluate(X_test, y_test, verbose=2)
```

```
814/814 - 3s - loss: 0.5946 - accuracy: 0.8239
```

## ▼ 4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the model with maximum probability.

```
import os
print(os.getcwd())
```

```
/content
```


```
checkpoint_best_path = 'checkpoints_best_only/checkpoint'
model_MLP = get_model(X_train[0].shape)
model_MLP.load_weights(checkpoint_best_path)
```

```
↳ <tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x7f4e8d326da0>
```

```
checkpoints_best_path = 'checkpoints_best_only_CNN/checkpoint'  
model_CNN = get_new_model(X_train[0].shape)  
model_CNN.load_weights(checkpoints_best_path)
```

```
↳ <tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x7f4e8d1de8d0>
```

```
import matplotlib.pyplot as plt  
  
fig, ax = plt.subplots(1, 5, figsize=(5, 1))  
for i in range(5):  
    ax[i].set_axis_off()  
    ax[i].imshow(X_test[i,:,:,:0])
```



```
num_test_images = X_test.shape[0]  
  
random_inx = np.random.choice(num_test_images, 5)  
random_test_images = X_test[random_inx, ...]  
random_test_labels = y_test[random_inx, ...]  
  
predictions = model_MLP.predict(random_test_images)  
  
fig, axes = plt.subplots(5, 2, figsize=(16, 12))  
fig.subplots_adjust(hspace=0.4, wspace=-0.2)  
  
for i, (prediction, image, label) in enumerate(zip(predictions, random_test_images, random_test_labels)):  
    axes[i, 0].imshow(np.squeeze(image))  
    axes[i, 0].get_xaxis().set_visible(False)  
    axes[i, 0].get_yaxis().set_visible(False)
```

```
axes[i, 0].text(10., -1.5, f'Digit {label}')
```

```
axes[i, 1].bar(np.arange(len(prediction)), prediction)
```


```
axes[i, 1].set_xticks(np.arange(len(prediction)))
```

```
axes[i, 1].set_title(f"Categorical distribution. Model prediction: {np.argmax(prediction)}")
```

```
plt.show()
```





```
num_test_images = X_test.shape[0]
```

```
random_inx = np.random.choice(num_test_images, 5)
```

```
random_test_images = X_test[random_inx, ...]
```

```
random_test_labels = y_test[random_inx, ...]
```

```
predictions = model_CNN.predict(random_test_images)
```

```
fig, axes = plt.subplots(5, 2, figsize=(16, 12))
```

```
fig.subplots_adjust(hspace=0.4, wspace=-0.2)
```

```
for i, (prediction, image, label) in enumerate(zip(predictions, random_test_images, random_test_labels)):
    axes[i, 0].imshow(np.squeeze(image))
    axes[i, 0].get_xaxis().set_visible(False)
    axes[i, 0].get_yaxis().set_visible(False)
    axes[i, 0].text(10., -1.5, f'Digit {label}')
    axes[i, 1].bar(np.arange(len(prediction)), prediction)
    axes[i, 1].set_xticks(np.arange(len(prediction)))
    axes[i, 1].set_title(f"Categorical distribution. Model prediction: {np.argmax(prediction)}")
```

```
plt.show()
```



